

Parallel Programming HW3 report

Blocked floyd warshall

蔡茗鈞
110030014

Implementation

- a. Which algorithm do you choose in hw3-1?
sequential version of blocked floyd warshall.
- b. How do you divide your data in hw3-2, hw3-3?
將dist matrix以block size 32×32 分割成數個block, 分給kernel的block執行
若n不為32的倍數, 將n補至32的倍數, 並在多出來的部分填入INF(padding)。
- c. What's your configuration in hw3-2, hw3-3? And why? (e.g. blocking factor, #blocks, #threads)
blocking factor = 32, 因為kernel的max thread per block為1024(32×32)
 $\#block = (n/32)^2$ (此時n為32的倍數)
 $\#threads = 32 \times 32$ (per block)
- d. How do you implement the communication in hw3-3?
一開始兩個gpu先各自利用cudaMemcpy(, , cudaMemcpyHostToDevice);
取得上半部與下半部的dist matrix, 進入round的for loop之後的每個iteration一開始, 利用cudaMemcpyPeer間gpu之間傳送, 各自取得剩下一半的dist matrix再開始計算, 因為iteration r, phase3只會計算一半的blocks (gpu 0 算上半部, gpu1計算下半部)。
- e. Briefly describe your implementations in diagrams, figures or sentences.
n = padding後的vertex數
B = 32
round = n/B
每個round r分成phase1, phase2 phase3。
phase1 計算block(r,r) (1個block)

phase2計算除了block(r,r)以外, row = r或column= r的所有block($2 \times (\text{round}-2)$ 個block)

phase3計算剩餘的所有block($(\text{round}-1)^2$ 個block)

每個phase的每個block將各自的dependency blocks從global memory複製進per block shared memory.

phase1的dependency block只有block(r,r)

phase2的dependency block有自身的block(i,j)與block(r,r)

phase3的dependency block有自身的block(i,j)與block(r,j)和block(i,r)

Profiling Results (hw3-2)

on testcase c21.3, n=5000 , kernel = phase3

a. global load/ store throughput (Min Max Avg)

Global Load Throughput 15.607GB/s 15.739GB/s 15.726GB/s

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	gld_throughput	Global Load Throughput	304.41MB/s	344.83MB/s	319.46MB/s
Kernel: phase3(int*, int, int, int)					
157	gld_throughput	Global Load Throughput	15.607GB/s	15.739GB/s	15.726GB/s
Kernel: phase2(int*, int, int, int, bool)					
314	gld_throughput	Global Load Throughput	12.068GB/s	13.398GB/s	13.072GB/s

Global Store Throughput 24.590GB/s 24.677GB/s 24.651GB/s

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	gst_throughput	Global Store Throughput	544.96MB/s	589.71MB/s	583.64MB/s
Kernel: phase3(int*, int, int, int)					
157	gst_throughput	Global Store Throughput	24.590GB/s	24.677GB/s	24.651GB/s
Kernel: phase2(int*, int, int, int, bool)					
314	gst_throughput	Global Store Throughput	19.889GB/s	22.733GB/s	22.129GB/s

b. shared memory load/ store throughput (Min Max Avg)

Shared Memory Load Throughput 1578.3GB/s 1585.9GB/s 1584.7GB/s

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	shared_load_throughput	Shared Memory Load Throughput	51.090GB/s	55.825GB/s	54.865GB/s
Kernel: phase3(int*, int, int, int)					
157	shared_load_throughput	Shared Memory Load Throughput	1578.3GB/s	1585.9GB/s	1584.7GB/s
Kernel: phase2(int*, int, int, int, bool)					
314	shared_load_throughput	Shared Memory Load Throughput	1087.9GB/s	1462.1GB/s	1419.6GB/s
n = 5000					

Shared Memory Store Throughput 49.272GB/s 54.807GB/s 52.072GB/s

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	shared_store_throughput	Shared Memory Store Throughput	8.4238GB/s	21.033GB/s	19.499GB/s
Kernel: phase3(int*, int, int, int)					
157	shared_store_throughput	Shared Memory Store Throughput	49.272GB/s	54.807GB/s	52.072GB/s
Kernel: phase2(int*, int, int, int, bool)					
314	shared_store_throughput	Shared Memory Store Throughput	38.582GB/s	50.950GB/s	46.546GB/s
n = 5000					

c. occupancy (Min Max Avg)

Achieved Occupancy 0.975517 0.975901 0.975632

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	achieved_occupancy	Achieved Occupancy	0.497012	0.497279	0.497118
Kernel: phase3(int*, int, int, int)					
157	achieved_occupancy	Achieved Occupancy	0.975517	0.975901	0.975632
Kernel: phase2(int*, int, int, int, bool)					
314	achieved_occupancy	Achieved Occupancy	0.932410	0.947622	0.938690
n = 5000					

d. sm efficiency (Min Max Avg)

Multiprocessor Activity 99.62% 99.92% 99.89%

Invocations	Metric Name	Metric Description	Min	Max	Avg
Device "NVIDIA GeForce GTX 1080 (0)"					
Kernel: phase1(int*, int, int)					
157	sm_efficiency	Multiprocessor Activity	3.41%	3.63%	3.57%
Kernel: phase3(int*, int, int, int)					
157	sm_efficiency	Multiprocessor Activity	99.62%	99.92%	99.89%
Kernel: phase2(int*, int, int, int, bool)					
314	sm_efficiency	Multiprocessor Activity	88.11%	89.96%	89.28%

Experiment & Analysis(hw3-2)

on testcase c21.3, n=5000

a. System Spec: Hades

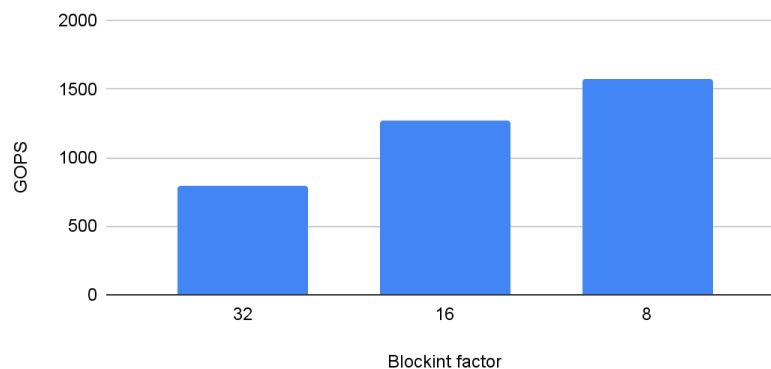
b. Blocking Factor (hw3-2)

GOPS

利用--metric instr_integer取得每個kernel得平均instruction count

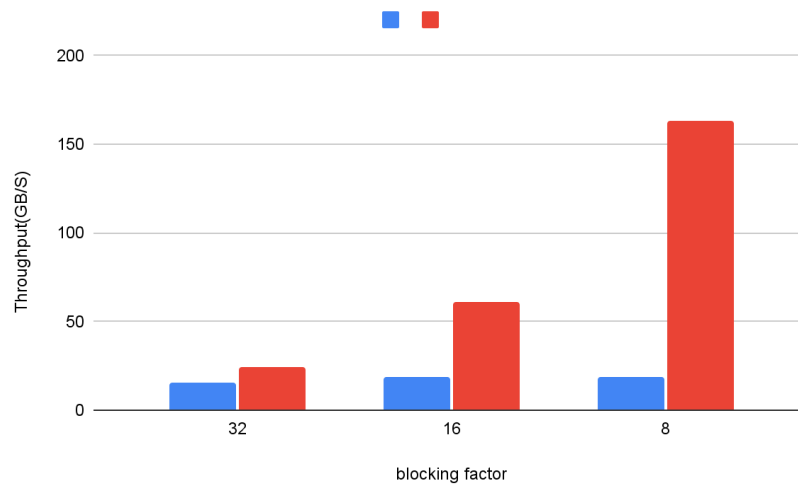
然後計算 $GOPS =$

$$(\#phase1_instr * \#phase1_calls + \#phase2_instr * \#phase2_calls + \#phase3_instr * \#phase3_calls) / (\#phase1_totaltime + \#phase2_totaltime + \#phase3_totaltime)$$



Global memory throughput (kernel = phase3)

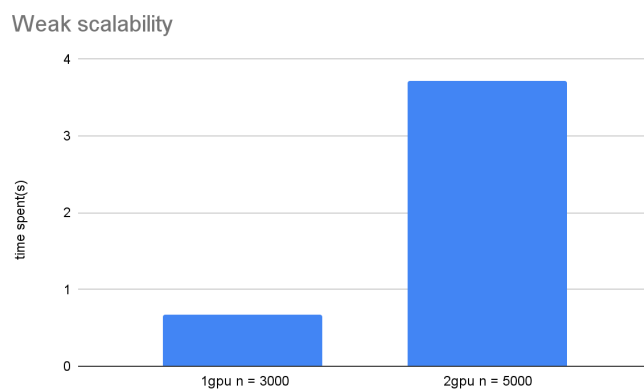
紅色:global memory store 藍色:global memory load



c. Optimization(hw3-2)

- shared memory:
for each block copy its dependency block to per block shared memory
- Reduce communication time:
use `cudaHostRegister(Dist, n*n*sizeof(int), cudaHostRegisterDefault);`
to pinned host memory for faster memcopy
- Handle bank conflict:
use `int j = threadIdx.x; int i = threadIdx.y;` instead of
`int j = threadIdx.y; int i = threadIdx.x;` to reduce bank conflict

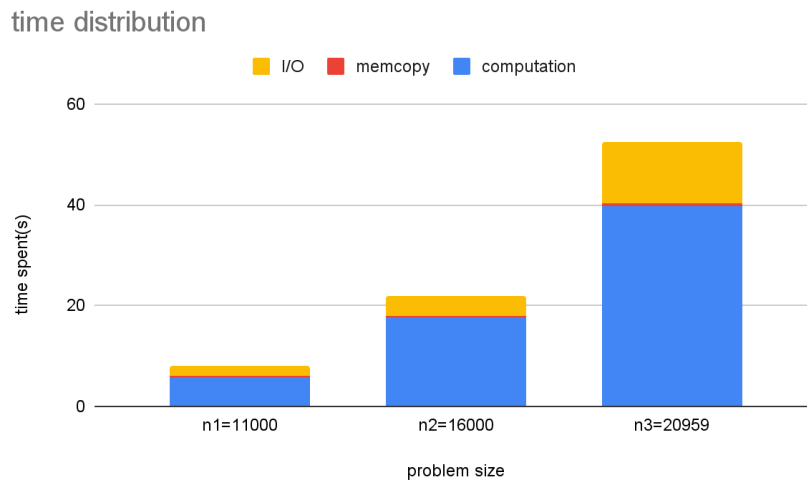
d. Weak Scalability(hw3-3)



$$(5000/3000)^2 \approx 2$$

the time significantly increased when i used 2 gpus

e. Time Distribution(hw3-2)



computation is the main bottleneck for my program

Experience & conclusion

Through this homework I learnt how to better utilize gpu to compute these kinds of computation. Although I didn't manage to perfect the result, I have some ideas to improve. For example, although I set the thread per block to 1024, which is the maximum number to set, I didn't fully use the shared memory, which is the first thing I would try to improve if I had to.

